Ethical Analysis of Artificial Intelligence: Societal Impacts and Moral Frameworks

Group: D
Members: Avila Concepción, Sergio; Chamorro Messa, Sofia; Muñiz Ríos, Andrés; Luis A Hernandez;
Course: ICOM5015 – Artificial Intelligence
Date: May 29, 2026

Content Table **Pages**

1. Introduction

Artificial Intelligence (AI) has evolved rapidly from a specialized research field into a technology that directly influences daily life, industry, education, politics, and scientific research. Systems based on deep learning and neural networks, such as ChatGPT, DALL·E, Google Gemini, and IBM Watson, demonstrate capabilities that once seemed impossible. However, this rapid progress has also generated ethical concerns related to labor displacement, surveillance, misinformation, environmental sustainability, and global inequality.

This report analyzes AI development through the perspectives of consequentialism, deontological ethics, and virtue ethics. Rather than treating AI only as a technical innovation, the report evaluates how these technologies transform social structures and human decision-making. The analysis focuses on labor markets, environmental impact, global inequality, misinformation, healthcare and education, and privacy rights.

The group organized the project by dividing responsibilities among members according to their interests and strengths. Research was conducted using academic articles, institutional reports, and reliable technological sources. The final objective is not to classify AI as entirely "good" or "bad," but to examine the conditions under which it can either improve human well-being or intensify social and economic problems.

2. Societal Impacts of Artificial Intelligence

AI generates both major benefits and important risks depending on how it is designed and implemented. In many sectors, it improves efficiency, accelerates scientific discovery, and expands access to information. Nevertheless, these advances are often accompanied by ethical concerns involving inequality, transparency, and concentration of power.

One of the most visible impacts of AI appears in employment and labor markets. While automation traditionally replaced repetitive manual work, generative AI now automates cognitive and creative tasks as well. Systems capable of generating software code, legal summaries, technical reports, and digital content increase productivity, but they also threaten administrative and creative professions, creating concerns about unemployment and wage inequality [1], [8].

Environmental impact has also become an important issue. Large AI models require enormous computational resources, generating high energy consumption, carbon emissions, and water usage for cooling data centers [2], [10]. Despite these costs, AI also contributes positively to climate research, energy optimization, and sustainability modeling [3].

Another concern is the concentration of technological power. Developing advanced AI systems requires specialized hardware, large datasets, and substantial financial investment. Consequently, a small number of corporations and developed nations dominate the AI ecosystem, increasing technological dependence among developing countries [11].

AI additionally creates opportunities for malicious misuse. Deepfakes and synthetic media can be used for fraud, political manipulation, impersonation, and reputational attacks [5]. As these systems become more realistic, distinguishing authentic content from fabricated information becomes increasingly difficult.

Despite these concerns, AI has produced major advances in healthcare, education, and scientific research. Machine learning assists medical diagnosis, supports personalized education, and accelerates biological research. Systems such as AlphaFold have significantly advanced protein structure prediction, demonstrating AI's potential to improve quality of life globally [15].

Finally, AI-driven surveillance systems raise concerns about privacy and civil liberties. Facial recognition, behavioral tracking, and large-scale data collection increase the ability of governments and corporations to monitor individuals, potentially threatening autonomy and freedom of expression.

3. Ethical Analysis

Employment and Labor Markets

The ethical debate surrounding automation centers on the balance between economic efficiency and human well-being. Consequentialism may justify AI automation if it increases productivity, lowers costs, and benefits society overall. However, it must also consider the negative effects of unemployment, economic insecurity, and social instability caused by labor displacement.

Deontological ethics argues that workers should not be treated merely as instruments for profit. Replacing employees without retraining opportunities or fair transition mechanisms disregards their dignity and autonomy. Corporations therefore have ethical responsibilities toward workers beyond economic efficiency.

Virtue ethics focuses on the moral character of institutions. Ethical organizations should implement technological changes responsibly, balancing innovation with social responsibility and employee development. The group concluded that consequentialism is useful for evaluating labor-market changes, but it must be limited by deontological protections that safeguard workers' rights.

Environmental Consequences

AI systems consume significant quantities of energy and water. From a consequentialist perspective, these environmental costs may be justified when AI contributes to scientific progress, climate modeling, or energy optimization [3]. However, some AI applications produce only limited social value while requiring extensive computational resources.

Deontological ethics frames sustainability as a moral duty. Organizations should operate transparently regarding emissions and resource consumption instead of concealing environmental costs. Virtue ethics emphasizes moderation, stewardship, and long-term responsibility. For this reason, the group considered virtue ethics the most appropriate framework for evaluating environmental impacts.

Global Inequalities and Technological Barriers

The concentration of AI infrastructure among a few corporations and technologically advanced nations raises concerns about global inequality. Consequentialism criticizes this concentration because it limits the distribution of technological benefits.

Deontological ethics provides a stronger critique by emphasizing fairness and justice. When advanced AI systems remain inaccessible because of financial or technological barriers, developing nations become dependent consumers rather than active participants in innovation. Virtue ethics promotes solidarity, cooperation, and knowledge sharing. The group concluded that deontological ethics best addresses issues of global inequality because these problems fundamentally concern rights and justice.

Deepfakes and Misinformation

Generative AI has significantly reduced the cost of producing deceptive content. Deepfakes and synthetic media can manipulate political discourse, facilitate fraud, and damage public trust [5]. Consequentialism evaluates these systems according to their social consequences, recognizing that misinformation can seriously weaken trust in digital information.

Deontological ethics condemns deepfakes because deception violates individual autonomy and dignity. Virtue ethics emphasizes honesty and integrity, arguing that ethical developers should implement safeguards such as watermarking and verification systems before releasing powerful generative tools. The group considered deontological ethics the clearest framework for regulating misinformation because truthfulness and consent should not be negotiable principles.

Healthcare, Education, and Scientific Research

AI applications in healthcare, education, and science offer significant social benefits. Consequentialism strongly supports these technologies because they can improve quality of life, reduce suffering, and expand educational opportunities.

However, deontological ethics introduces important limits. Patients and users have rights to transparency and informed consent, especially when black-box systems are involved. Human oversight remains essential in sensitive fields such as healthcare. Virtue ethics highlights the importance of empathy and professional judgment, arguing that AI should complement rather than replace human expertise.

Privacy and Civil Liberties

Modern AI systems rely heavily on surveillance, behavioral analysis, and large-scale data collection. Consequentialists may justify these systems in the name of security or efficiency, but such arguments often ignore the long-term social costs of constant monitoring.

Deontological ethics rejects mass surveillance because it violates privacy and individual autonomy. Personal data should not be collected without informed consent. Virtue ethics questions the morality of institutions that depend on excessive surveillance, arguing that societies based on constant monitoring promote distrust rather than mutual respect. The group concluded that deontological ethics provides the strongest framework for protecting privacy and civil liberties.

4. Group Deliberation

During the final meeting, the group discussed several disagreements between ethical frameworks. One of the main debates involved surveillance technologies. Some members argued from a utilitarian perspective that AI-based surveillance systems can reduce crime and improve public safety. Others maintained that even effective surveillance systems remain ethically unacceptable if they violate individual rights and autonomy. After discussion, the group concluded that fundamental rights should establish limits on utilitarian calculations.

Another debate emerged around the use of black-box AI systems in healthcare. Consequentialist arguments emphasized that highly accurate diagnostic systems could save lives even if their reasoning processes remain opaque. In contrast, deontological concerns focused on transparency and accountability. The group ultimately agreed that AI should support medical professionals rather than replace human judgment entirely.

These discussions demonstrated that no single ethical theory adequately resolves all AI-related dilemmas. As a result, the group adopted a pluralistic position that combines consequentialist evaluation of large-scale outcomes, deontological protection of rights, and virtue ethics as guidance for professional conduct and institutional culture.

5. Conclusion

Artificial Intelligence represents one of the most transformative technologies of the modern era. Its applications have the potential to improve healthcare, education, scientific discovery, and economic productivity. At the same time, AI systems can intensify inequality, weaken privacy protections, accelerate environmental damage, and destabilize democratic institutions if deployed irresponsibly.

This report shows that ethical evaluation of AI requires more than technical analysis. The central challenge is determining how technological innovation can be aligned with human dignity, social justice, and long-term societal well-being. Consequentialism helps evaluate broad social outcomes, deontological ethics establishes limits that protect fundamental rights, and virtue ethics encourages the development of responsible institutional cultures.

Our group concludes that AI should not be treated as an autonomous force beyond human control. The future impact of these systems will depend on the ethical standards enforced by governments, corporations, researchers, and society itself. Responsible governance, transparency, international cooperation, and ethical education will therefore play a decisive role in shaping whether AI becomes a tool for collective flourishing or a source of deeper global instability.

Appendix

Minutes of the First Meeting

Course: ICOM5015 – Artificial Intelligence
Group: D
Date: May 22, 2026
Topic: Organization and planning of the project "Ethical Analysis of Artificial Intelligence: Societal Impacts and Moral Frameworks"

Development of the meeting

During the first meeting, the group members discussed the assignment requirements and analyzed the overall approach the report would take. It was agreed that the work would focus on the ethical analysis of Artificial Intelligence and its impacts on society, based on three main philosophical currents: consequentialism, deontological ethics, and virtue ethics.

The group identified the main areas that needed to be investigated to develop the analysis. Among the selected topics were:

- Impact of AI on employment and the labor market.
- Environmental consequences.
- Global technological inequality.
- Disinformation and deepfakes.
- Applications in health and education.
- Privacy and surveillance.

Subsequently, responsibilities were distributed among the group members, taking into account each member's interests and strengths. It was also agreed to use academic articles, institutional reports, and other reliable sources to support the report's content.

Decisions made

1. The report will be developed from an ethical and social perspective, not just a technical one.
2. The ethical theories of the following will be used:
    - Consequentialism.

    - Deontological ethics.
    - Virtue ethics.
3. The report sections will be divided among the team members to facilitate the research.
4. Academic references and recognized sources will be used to support the analysis.
5. A final meeting will be scheduled to integrate the results and discuss the conclusions of the work.

Minutes of the Last Meeting

Course: ICOM5015 – Artificial Intelligence
Group: D
Date: May 26, 2026
Topic: Final Integration and Ethical Discussion of the Project

Meeting Outline

During the last meeting, the group members presented the results obtained in their respective research projects, and all sections of the final report were integrated.

The group discussed different viewpoints related to the ethical dilemmas of Artificial Intelligence. One of the main debates focused on the use of AI-based surveillance systems. Some members defended the idea that these technologies can improve public safety and reduce crime; however, others argued that such practices can violate privacy and fundamental human rights.

The use of "black-box" AI systems in healthcare was also discussed. The benefits of using highly accurate algorithms for medical diagnoses were analyzed, but the importance of maintaining human oversight, transparency, and accountability in decision-making was also emphasized.

Finally, the group reviewed the overall wording of the report, verified the bibliographic references, and finalized the document's structure.

Decisions Made

1. It was concluded that no single ethical theory can fully resolve all the problems associated with AI.
2. The group adopted a pluralistic position that combines:

- evaluation of social consequences,
- protection of fundamental rights, and
- moral and institutional responsibility.

3. It was agreed that AI should be used as a support tool and not as a complete replacement for human judgment, especially in critical areas such as medicine.
4. It was established that human rights and privacy must act as limits against the excessive use of surveillance systems.
5. The final version of the report was approved for submission.

Individual reflections

- During the course of this project, I learned that Artificial Intelligence should not only be analyzed from a technological perspective, but also from the perspective of its social and ethical consequences. I understood how the use of AI can benefit areas such as medicine and education, but also generate problems related to privacy, unemployment, and misinformation. My contribution to the group consisted of researching academic information on the social impacts of AI and collaborating on organizing the main ideas of the report. (Sofia Chamorro)
- This project allowed me to better understand the importance of ethical theories, especially consequentialism, deontological ethics, and virtue ethics, for evaluating modern technologies like AI. I learned that technological decisions must consider both the benefits and the rights of individuals. During the project, I contributed by conducting research related to surveillance, privacy, and the impact of AI systems on society. (Luis A Hernández)
- Through this activity, I learned that technological advancement also entails social and professional responsibility. I was particularly struck by how AI can exacerbate inequalities between countries and concentrate technological power in the hands of a few companies. My contribution to the group was participating in the ethical discussions and assisting in the drafting and review of the report's final conclusions. (Andrés E. Muñiz Ríos)
- This work strengthened my ability to evaluate complex problems from multiple ethical perspectives, rather than relying solely on technical criteria. I learned that the challenges posed by Artificial Intelligence cannot be adequately addressed from a single ethical perspective, making it necessary to integrate consequentialist, deontological, and virtue-based approaches to underpin responsible decision-making.

My main contribution consisted of drafting and coordinating the integration of the report's sections, refining the coherence of the arguments, and ensuring clarity and consistency throughout the final document.(Sergio Avila C.)

References

[1] J. Hatzius et al., “The Potentially Large Effects of Artificial Intelligence on Economic Growth,” Goldman Sachs Global Economics Analyst, 2023.

[2] E. Strubell, A. Ganesh, and A. McCallum, “Energy and policy considerations for deep learning in NLP,” in *Proceedings of ACL 2019*, Florence, Italy, 2019, pp. 3645–3650.

[3] D. Rolnick et al., “Tackling climate change with machine learning,” *ACM Computing Surveys*, vol. 55, no. 2, pp. 1–96, 2022.

[4] C. Cath, “Governing artificial intelligence: ethical, legal and technical opportunities and challenges,” *Philosophical Transactions of the Royal Society A*, vol. 376, no. 2133, 2018.

[5] M. Brundage et al., “The Malicious Use of Artificial Intelligence: Forecasting, Prevention, and Mitigation,” Future of Humanity Institute, University of Oxford, 2018.

[6] I. Goodfellow et al., “Generative adversarial networks,” in *Proceedings of NIPS 2014*, Montreal, Canada, 2014, pp. 2672–2680.

[7] J. Manyika, J. Silberg, and B. Presten, “What do we do about the biases in AI?” *Harvard Business Review*, Oct. 2019.

[8] World Economic Forum, “The Future of Jobs Report 2024,” Geneva, Switzerland, 2024.

[9] S. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*, 4th ed. Pearson, 2021.

[10] P. Li et al., “Making AI Less ‘Thirsty’: Uncovering and Addressing the Secret Water Footprint of Large Language Models,” arXiv preprint arXiv:2304.03271, 2023.

[11] K. Crawford, *Atlas of AI: Power, Politics, and the Planetary Costs of Artificial Intelligence*. Yale University Press, 2021.

[12] J. Rawls, *A Theory of Justice*. Harvard University Press, 1971.

[13] European Parliament and Council, “Regulation (EU) 2016/679 (GDPR),” *Official Journal of the European Union*, 2016.

[14] I. Kant, *Groundwork of the Metaphysics of Morals*. Cambridge University Press, 1998.

[15] J. Jumper et al., “Highly accurate protein structure prediction with AlphaFold,” *Nature*, vol. 596, no. 7873, pp. 583–589, 2021.

[16] Aristotle, *Nicomachean Ethics*. Cambridge University Press, 2000.

[17] N. Bostrom, *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press, 2014.

[18] S. Zuboff, *The Age of Surveillance Capitalism*. PublicAffairs, 2019.